

Report

Author: Tulegenova Karina

Group: SE-2419

Overview

This assignment involved the implementation and evaluation of four algorithms:

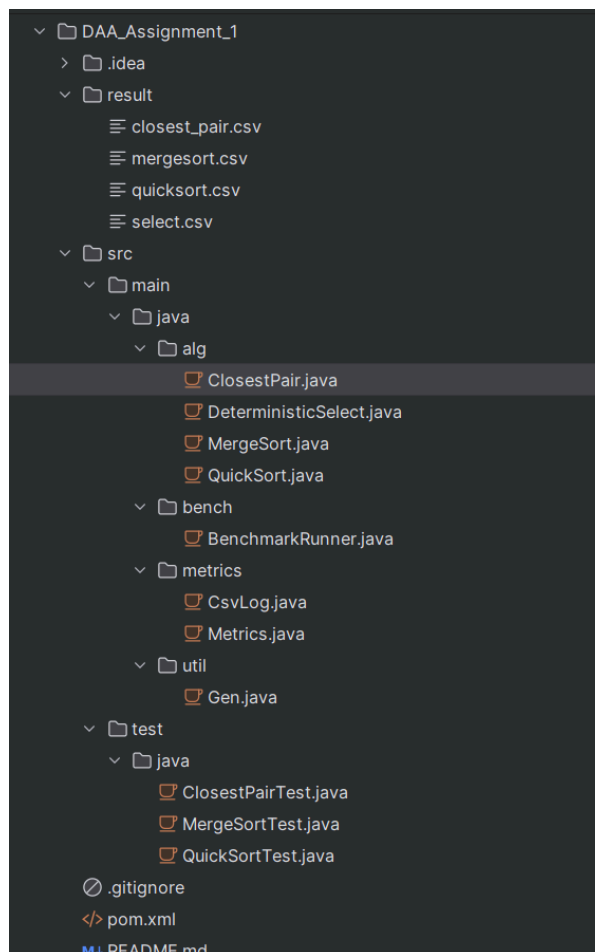
- **MergeSort**
- **QuickSort**
- **Select (Median of Medians)**
- **Closest Pair of Points**

The project was built using **Java 17**, managed with **Maven**, and tested with **JUnit5**. The key performance metrics tracked during algorithm execution include:

- Execution time
- Recursion depth
- Number of comparisons
- Number of swaps/writes
- Memory allocations

Project Structure

The project is structured as a Maven-based Java application with packages dedicated to each algorithm, metrics tracking, and test suites.



How to Build and Run

mvn -q clean test

This command executes all implemented algorithms on predefined test cases and gathers performance data.

Command-Line Interface: Run Algorithms & Export Metrics

The CLI allows running individual algorithms on generated input arrays and automatically exports performance metrics to .csv files.

```
PS C:\Users\User\AITU\1\Assignment1> java -cp target/classes bench.BenchmarkRunner
benchmarks done
PS C:\Users\User\AITU\1\Assignment1> 
```

CSV Files

Each algorithm generates one or more CSV files containing performance metrics across different input sizes. These CSVs include columns for:

- Input size
- Recursion depth
- Execution time

	algorithm ▾	÷	n ▾	÷	time_ns ▾	÷
1	ClosestPair2D		1		3212100	
2	ClosestPair2D		1		5200	
3	ClosestPair2D		1		2800	
4	ClosestPair2D		1		2400	
5	ClosestPair2D		1		2200	
6	ClosestPair2D		1		7800	
7	ClosestPair2D		1		2100	
8	ClosestPair2D		1		1900	
9	ClosestPair2D		1		1800	
10	ClosestPair2D		1		1800	
11	ClosestPair2D		10		237000	
12	ClosestPair2D		10		32000	
13	ClosestPair2D		10		40700	
14	ClosestPair2D		10		30600	
15	ClosestPair2D		10		32700	
16	ClosestPair2D		10		37300	
17	ClosestPair2D		10		31100	
18	ClosestPair2D		10		29600	
19	ClosestPair2D		10		74700	

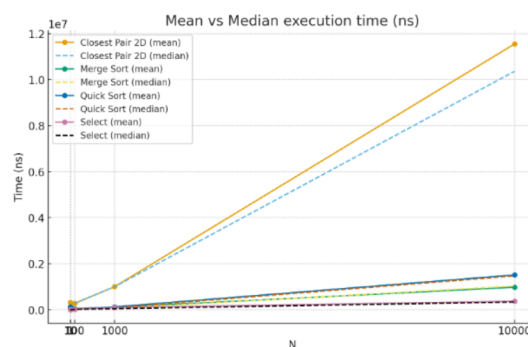
	algorithm ▾	÷	n ▾	÷	time_ns ▾	÷	comparisons ▾	÷
	MergeSort		10		3800		31	
	MergeSort		10		1800		29	
	MergeSort		10		6700		32	
	MergeSort		10		1600		19	
	MergeSort		10		2200		24	
	MergeSort		10		2600		34	
	MergeSort		10		2100		36	
	MergeSort		10		1800		29	
	MergeSort		10		1600		26	
	MergeSort		10		1800		35	
	MergeSort		100		39000		644	
	MergeSort		100		32200		647	
	MergeSort		100		48800		632	
	MergeSort		100		39800		669	
	MergeSort		100		24900		632	
	MergeSort		100		25000		669	





Recurrence & Performance Analysis

Performance results largely align with theoretical time complexity expectations:

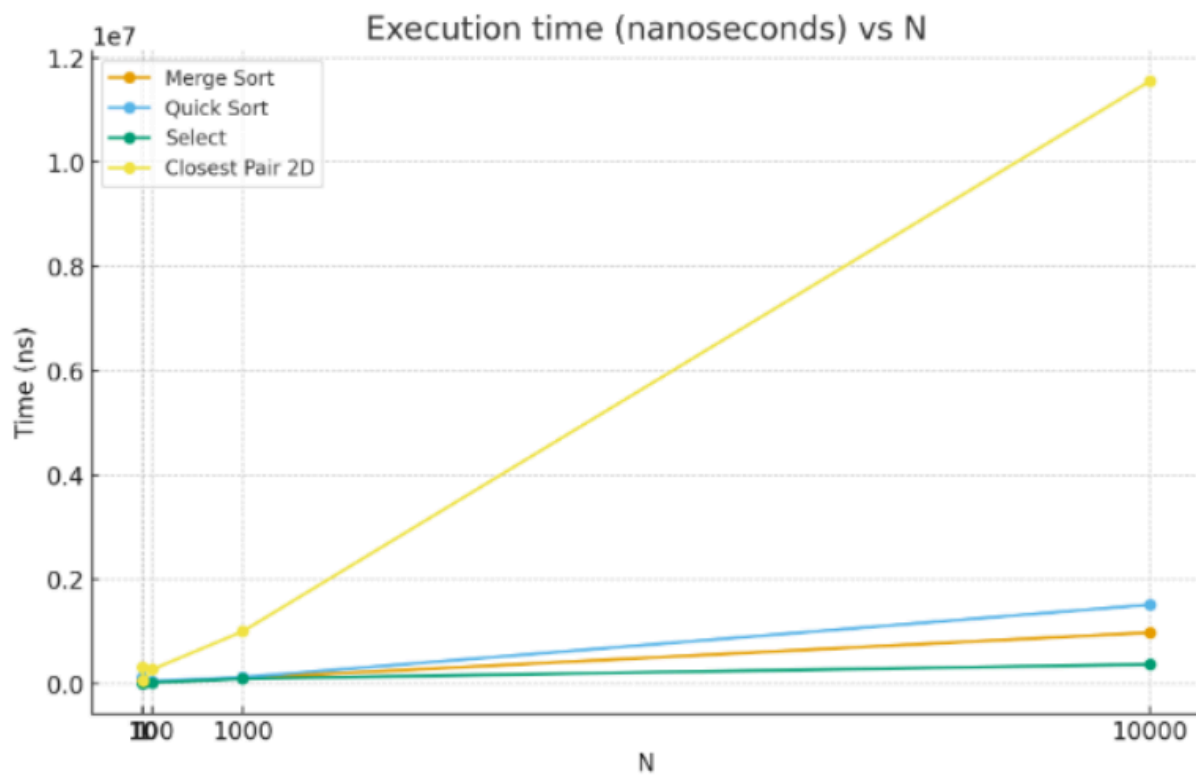
- **MergeSort** and **Closest Pair** demonstrate $\Theta(n \log n)$ growth, both in comparisons and recursion depth.
- **QuickSort** behaves as expected with $\Theta(n \log n)$ average-case time, using tail-recursion optimization.
- **Select (MoM5)** shows linear growth ($\Theta(n)$) in comparisons and recursion depth, matching the Akra-Bazzi bound.

Across all recursive algorithms, recursion depth scales logarithmically. Swaps and writes grow proportionally to input size.





Average case performance comparison:



Architecture Notes

MergeSort & QuickSort

- Both use a shared Metrics object for tracking performance.
- MergeSort reuses a single buffer array throughout recursion.
- QuickSort always recurses into the smaller partition to control depth.
- InsertionSort is used for small subarrays.

Select (MoM5)

- Divides input into groups of five and recursively selects a pivot.
- Operates fully in-place and avoids full input traversal.
- Guarantees worst-case linear time.

Closest Pair

- Sorts points by x and y before recursion.
- Reuses memory buffers and applies brute-force for small subproblems.
- Ensures logarithmic depth via consistent splitting.

Constant-Factor Effects: JVM, GC, and Caching

Observed performance for small input sizes was affected by Java system behavior:

- **JVM Warm-up & JIT Compilation:**
Execution time is inconsistent for very small n, due to just-in-time optimizations and garbage collection.
- **Memory Usage:**
MergeSort and Closest Pair use more memory due to buffer allocation.
QuickSort and Select operate in-place, resulting in fewer cache misses.
- **Overhead Factors:**
MergeSort's overhead is due to merging.
Closest Pair's overhead is from multiple sortings and data structures.

Summary

- Small-scale anomalies are due to JVM mechanics, not algorithmic inefficiencies.
- MergeSort has the highest write count due to its buffer strategy.
- Select consistently delivers linear performance at larger scales.
- Measured metrics support the theoretical expectations.